



```

Config: Batch=64 | Optimizer=SGD | Momentum=0.9 | weight decay: True | Margin=1.0 | LR: 5.0e-02 -> 5.0e-02 -> 5.0e-02 -> 5.0e-02 |
          total dataset size: 10000 |
          train transformations: Compose(
            Resize(size=(224, 224)), interpolation=Bilinear, max_size=None, antialias=True)
            RandomPerspective(p=0.3)
            RandomHorizontalFlip(p=0.3)
            RandomRotation(degrees=[-15.0, 15.0], interpolation=nearest, expand=False, fill=0)
            RandomGrayscale(p=0.1)
            ColorJitter(brightness=[0.8, 1.2], contrast=[0.8, 1.2], saturation=[0.8, 1.2], hue=[-0.1, 0.1])
            ToTensor()
            Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
          )
          Emb=Sequential(
            (0): Linear(in_features=2048, out_features=1024, bias=True)
            (1): ReLU()
            (2): Linear(in_features=1024, out_features=128, bias=True)
          )
          freeze: 5 -> 4 (at ep 5) -> 3 (at ep 15) -> 3 (at ep 22)
          LR decrease factor: 2
    
```