



```
Config: Batch=64 | With Normalization on output | With lower backbone LR by a factor of 0.5 | Optimizer= SGD | weight decay: 0 | Margin=0.2 | LR: 1.0e-02 -> 1.0e-02 |
total dataset size: 10000 |
train transformations: Compose(
    Resize(size=(224, 224), interpolation=bilinear, max_size=None, antialias=True),
    RandomHorizontalFlip(p=0.3),
    RandomRotation(degrees=[-15.0, 15.0], interpolation=nearest, expand=False, fill=0),
    RandomGrayscale(p=0.1),
    ColorJitter(brightness=[0.8, 1.2], contrast=[0.8, 1.2], saturation=[0.8, 1.2], hue=[-0.1, 0.1]),
    ToTensor(),
    Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
)
Emb=Sequential(
    (0): Linear(in_features=2048, out_features=1024, bias=True)
        (1): ReLU()
    (2): Linear(in_features=1024, out_features=256, bias=True)
)
Freeze: 4 -> 3 (at ep 5)
LR decrease factor: 1
```