Loss Curve (cosine) — Train Loss, Val Loss

F1 Score Curve (cosine) — Val F1 Score

Threshold Evolution (cosine) — Best Threshold

Config: Batch=64 | With Normalization on output | With lower backbone LR by a factor of 0.5 | Optimizer= Adam | weight decay: 0 | Margin=0.2 | LR: 1.0e-04 -> 1.0e-05 |
total dataset size: 10000 |
train transformations: Compose(
Resize(size=(224, 224), interpolation=bilinear, max_size=None, antialias=True)
RandomHorizontalFlip(p=0.3)
RandomRotation(degrees=[-15.0, 15.0], interpolation=nearest, expand=False, fill=0)
RandomGrayscale(p=0.1)
ColorJitter(brightness=(0.8, 1.2), contrast=(0.8, 1.2), saturation=(0.8, 1.2), hue=(-0.1, 0.1))
ToTensor()
Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
) |
Emb=Sequential(
(0): Linear(in_features=2048, out_features=1024, bias=True)
(1): ReLU()
(2): Linear(in_features=1024, out_features=256, bias=True)
)
Freeze: 4 -> 3 (at ep 12)
LR decrease factoin: 0.1